KNOWLEDGE THAT COMPOUNDS

Watch once. Teach forever.

Drop in a **YouTube link**, a **web article**, a **PDF**, or a **doc**. Riverfed turns it into a **knowledge base**, then into **Marp decks**, **editable PowerPoint**, **Excalidraw diagrams**, and **speaker manuals**, all styled by a **design picker** you can edit or extend. Every dropped file makes the next one smarter.

Try the demo →
See how it works →

`.youtube` · `web article` · `.pdf` · `.docx` · `.pptx` · `.xlsx` · `.mp3` · `.mp4` · `.html` · `.epub` · `.md` · `.txt` · `+ more`

THE PROBLEM

You consume more than you remember. Riverfed turns it into an `asset`.

Every podcast, masterclass, and PDF you've watched in 2026 lives somewhere as a dead artefact: a tab you closed or a transcript you'll never re-read. Riverfed ingests it once and makes it compound: a wiki that gets denser, a deck that fans out into trainings, a vault your agents can read.

THE PIPELINE

Three steps. Fully automatic.

Add what you have. Organise it with the right knowledge model. Generate material people and agents can use.

STEP 01

Ingest

Drop a URL or drag in a file. We extract clean Markdown from transcripts, OCR'd PDFs, parsed slides, and transcribed audio. Each output includes frontmatter, deduplication, and source-URL tracking.

`→ output/*.md`

STEP 02

Organise

An agent reads the new source in the context of the whole base. The selected knowledge model decides the folders, relationships, promotion rules, and prompts.

`→ model folders + links + promotion rules`

STEP 03

Activate

The knowledge base feeds the Studio. Generate teaching decks, editable PowerPoint, diagrams, speaker guides, Anki cards, or agent skills according to the model.

`→ training · diagrams · flashcards · skills`

125	29	257	52
SOURCES CAPTURED	TOPIC PAGES	ENTITY PAGES	CONCEPT PAGES

Live from Wiki · Karpathy Wiki

THE SYSTEM

One engine. Many ways to organise knowledge.

The Karpathy teach-back wiki is one shipped model, not the whole product. Second Brain turns captured knowledge into projects and areas. SOPs-to-Skills turns operational knowledge into instructions agents can execute. Each model shares the same capture, governance, and output engine.

New agent practices can be added as model manifests. The Studio discovers them without a page or server rewrite.

Capture LAYER 1
Keep converted files and transcripts intact, traceable, and ready to reuse.

Organise LAYER 2
Shape each base with the model that fits the work: teach-back, PARA, operational skills, or a custom method.

Govern LAYER 3
Apply the model's schema, relationships, promotion thresholds, and agent instructions.

Activate LAYER 4
Turn the base into teaching material, decision support, spaced repetition, or executable agent skills.

KNOWLEDGE MODEL REGISTRY

Best practices become reusable models.

Each model packages its folders, relationships, promotion thresholds, ingest prompts, and generators. Add a model to the registry and it appears here automatically.

Karpathy Wiki

Three-layer teach-back knowledge base, immutable raw sources distilled into an LLM-curated wiki of sources, topics, entities, and concepts, governed by an explicit schema. The original Riverfed model.

FOLDERS
Sources · Topics · Entities · Concepts

OUTPUTS
Marp · Pptx · Diagram · Guide

Second Brain (PARA)

Tiago Forte's PARA method as a knowledge base, capture notes distilled toward ACTION, not teach-back. Sources are captured 1:1, then organized into Projects (active, outcome-driven), Areas (ongoing responsibilities), and Resources (reference interests), with Archives for what went inactive.

FOLDERS
Sources · Projects · Areas · Resources · Archives

OUTPUTS
Marp · Pptx · Diagram · Guide

SOPs to Skills

A small business's operational knowledge, dropped docs distilled into Procedures (step-by-step how-tos), Policies (rules and standards), and Runbooks (diagnose-and-recover recipes), structured so the Skills generator can emit a Claude Code skill library from it. Knowledge in, agents out.

FOLDERS
Sources · Procedures · Policies · Runbooks

OUTPUTS
Marp · Pptx · Diagram · Guide · Skills

WHAT YOU GET

One drop. A whole studio.

Every ingest emits a portable knowledge base plus the materials to teach it. The Studio keeps growing.

Model-shaped vault

The selected model defines the folders and links, from sources, topics, entities, and concepts to PARA or a skill library. Each base lives outside the product as its own portable, git-ready vault.

`~/RiverfedVaults/`

Auto-rendered Marp deck

~36 slides, brand-themed, real source citations, no bullet-soup. The Studio renders **HTML, PDF, and PPTX** the moment the agent writes the markdown. No manual Marp command.

`.md + .html + .pdf + .pptx`

Native editable PowerPoint

A sibling renderer builds a real .pptx in which every shape and text box is hand-editable. Choose from **13 slide types**: cover, lead, section, content, KPI, donut KPI, quote, versus, cards, six-up, flow, timeline, and hub.

`.pptx · 13 types`

Design picker + editor

Pick a palette + typography from the Studio, or build your own. Three ship by default (**Riverfed**, **Heritage**, **Aurora**) and the editor lets you customise colours + fonts. DESIGN.md spec.

`design.md`

Excalidraw diagrams

Brand-coloured whiteboard sketches: agent loops, architecture maps, comparisons. Drag the JSON onto excalidraw.com to edit.

`.excalidraw`

Smart frame ingest

For slide-deck or whiteboard videos, `--frames --dense-until` samples the first 20 min every 45 sec. The ingest agent reads each frame *as an image* and extracts structural diagrams (cycles, hierarchies, before/after flows) the transcript flattens to bullets.

`--frames`

Speaker manual + Word doc

Get a slide-by-slide script, rhetorical hooks designed to land in silence, anticipated audience Q&A, and pacing cues. Export as Markdown or a brand-styled .docx ready to share with stakeholders.

`.md + .docx`

Bring your own agent

Every pipeline step runs on the agent CLI you already pay for: **Claude Code** (claude.ai Pro/Max) or **Codex** (ChatGPT). Auto picks whichever is logged in; an explicit pick never silently switches. No API key.

`claude · codex`

Knowledge models

A base is a template, not a fixed shape. Choose the Karpathy-style teach-back wiki, a **Second Brain (PARA)** organized for action, or **SOPs-to-Skills** for operational docs. Tune thresholds + agent prompts in the browser.

`models/`

Watched channels

Subscribe a base to YouTube channels: new uploads convert by themselves (queue-for-approval by default; flip a channel to auto-ingest for fully unattended growth). Knowledge that compounds while you sleep.

`watcher`

Anki flashcards

Atomic, source-cited Q/A cards are authored from the base and rendered to a real `.apkg` with stable deck IDs. Regenerating *updates* your deck in Anki instead of duplicating it.

`.apkg`

Agent skill library

The SOPs-to-Skills compiler turns a base of procedures into an executable library with one `SKILL.md` per procedure, policies folded into gotchas, and an `AGENTS.md` file. It loads in Claude Code *and* Codex sessions.

`SKILL.md + AGENTS.md`

Audit trail

Every claim cites its source slug. Open the page, read the original transcript, verify in one click.

`log.md`

Re-runnable

Add a new transcript, re-run the generators: every artefact refreshes with new content. No manual editing.

`idempotent`

More generators on the way

Notion exports, podcast scripts, social-card carousels, handouts. Each new generator reads the same knowledge base. Write once, teach everywhere.

`coming soon`

INPUTS

Anything you can drop, we can read.

YouTube transcription, OCR for scanned PDFs, transcription for audio & video, full structure for Office and EPUB. One pipeline.

YT	YouTube URL	PDF	Document	DOCX	Word	PPTX	Slides
XLSX	Spreadsheet	EPUB	Book	MP3	Audio	WAV	Audio
MP4	Video	HTML	Web page	MD	Markdown	TXT	Plain text

QUESTIONS

What you might be wondering.

What does “automatic” actually mean? +

Does it work offline? +

Which AI subscription do I need? +

Is Karpathy the only knowledge model? +

Marp *or* PowerPoint? +

Can I customise the look? +

What about whiteboard / Excalidraw videos? +

What generators ship today, and what’s next? +

Who's it for? +

Start your knowledge base in 60 seconds.

Choose a model, drop a source, and turn it into connected knowledge your team, your training, and your agents can use.

Open the converter →